

ESE532 Project Final Report - Group

Ritika Gupta, Taylor Nelms, and Nishanth Shyamkumar

December 6, 2019

1 Introduction

The elephant in the room for the entirety of our implementation is that, as of 24 hours before the due date of this project, it is not functioning correctly. It gets through the entirety of the input, and (finally) does not hang in the middle of a deadlock or infinite loop somewhere in hardware, but upon encoding and decoding some files (particularly, long binary files), the decoded version does not match the original.

Now, that does not mean we cannot obtain reasonable ideas as to performance, area/time costs, etc. for our design. However, concerns such as specific design axes eluded us through development because we were unable to reach the essential milestone of "a working design."

This is all to say that, while we cannot empirically verify any results of design changes, or give good accounts of verification strategies, what we **can** do is evaluate design decisions we could have made through the process, and analyze what kind of effects they would have on a mythical end result where our software worked.

As such, you must excuse us if we are light on some empiric details, as we, even in this final stretch, endeavor to improve our design to the point of functionality. 90-minute build times are, unfortunately, not conducive to rapid prototyping, so we will do what we can to represent our control of the source material in an academic sense, if not always a practical one.

2 Single ARM processor design

2.1 Key parameters:

The key parameters involved in the original ARM processor design are:

Reading data from the Filesystem.

Processing fingerprints in Rabin CDC and passing the chunks to SHA.

SHA calculates the SHA on each chunk.

If there is no SHA match from the SHA table, it sends data to LZW.

LZW compresses the chunks.

Finally the chunk is packed into a file.

2.2 Performance achieved and Energy required:

The throughput performance achieved for an all software design, i.e, where each of the algorithms were running on the Arm core was 7Mbps.

The energy required for this design was $\text{Power} = 2.22\text{W}$

$\text{Time} = 40 \text{ minutes} = 2400 \text{ seconds}$

$\text{Energy} = \text{PT} = 2.22 * 2400$

$= 5328 \text{ Joules}$

2.3 Compression achieved:

Input file : vmlinuz.tar

Input size : 66MB

Compressed size: 36.6MB

2.4 Time spent on major components:

Rabin : 286011cycles
Throughput : 59Mbps

Sha : 407923 cycles
Throughput : 41.53Mbps

LZW : 1023284 cycles
Throughput : 16.56Mbps

3 Ultra 96 design

– Performance achieved and energy required – Compression achieved – Key design aspects: task decomposition, parallelism, mapping to Zynq resources, include diagrams to support – Be clear where each component of the final design is performed (e.g., ARM, NEON vector, FPGA logic). – Model to explain performance, area, and energy of design – Current bottleneck preventing higher performance

All the components CDC, SHA, LZW and deduplication operate in hardware. Input data is read from network and it is processed in 2MB parts.

3.1 Compression achieved:

Input file : LittlePrince.txt
Input size : 14KB
Compressed size: 4.4KB

Input file : ESE532.tar
Input size : 50KB
Compressed size: 24.2KB

Input file : Franklin.txt
Input size : 390KB
Compressed size: 308KB

Input file : vmlinuz.tar
Input size : 66MB
Compressed size: 36.6MB

Input file : linux.tar
Input size : 191MB
Compressed size: 60.9MB

3.2 Processor components:

Input data is read from network and it is processed in 2MB parts.

The input read over the network is handled using a pthread that is executing on its own core.

The primary processor core(core 0) is running the PS section that interfaces with the PL logic.

The network thread is bound to run only on core 2.

3.2.1 Network thread functionality

The thread uses a 200MB input buffer that is shared between the cores 2 and 0.

The NW thread then reads packets, calculates the length of each packet and then copies over the bytes to the buffer.

In our implementation, we use 2MB buffers to feed the dataflow model implemented on the FPGA.

So the NW thread has the additional responsibility to increment a counter everytime it crosses the 2MB boundary. This variable is checked against in the main processor to make certain that it does not try to process any data before the buffer data fills up. The variable is used as a synchronization mechanism.

Finally upon exit, the NW thread indicates that it has completed its purpose and marks a variable that is used by core 0 to figure out that no more packets are to be expected.

3.2.2 Output thread functionality

There is a 3rd thread outside of the main and network thread that is tasked with writing processed data onto a file.

The output thread is bound to core 3 and it is responsible for opening, closing files and of course writing the data to the file. It is invoked based on the existing synchronization mechanisms and it frees the processor 0 from having to deal with write backs to the file.

CDC performs chunking and the data is taken by SHA and LZW which operate parallelly. The output of SHA which is 256 bit SHA value and the compressed output of LZW are both taken by deduplicate section to finally write the appropriate output to the output file. 3 ARM cores are being used. One is used for reading the input data, one for processing the data and the third for writing the output to the output file.

3.3 FPGA Design

3.4 Performance achieved

For linux.tar, we got 193 Mbps. That is the maximum speed we got. For smaller files, the speed was less than this.

3.5 Energy measurements

Power = 2.861 W

Time = 16 s

Energy = 2.861 * 16 = 45.77

Power split among various components is as shown :

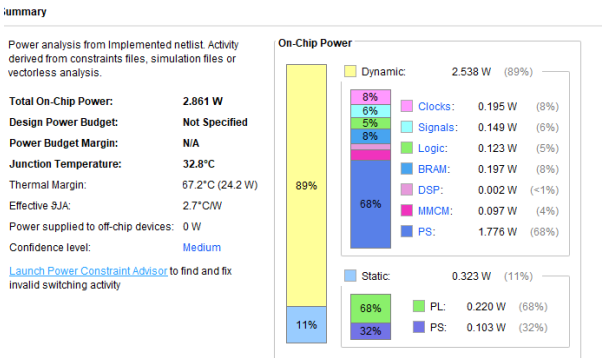


Figure 1: Power split.

3.5.1 Hardware Wrapper

On the outside of our PL functionality, we have a wrapper function that allows us to have a singular top-level function that operates on a larger buffer than a single chunk, but a smaller amount of data than an entire file.

For us, this has meant calling an overall function on a 2MB section of data. This was chosen to be larger than a single chunk by a few factors of scale, while being small enough to allow for sensible passing of data into the PL.

This wrapper function, `processBuffer`, sets up the relevant streams between components and calls each of those component functions. Each component is responsible for reading and writing from its relevant stream for as many times as it needs across the entire length of that input buffer.

It contains simple functions for turning the internal streams into the whole-buffer sizes that SDx likes, which allows us to use a `SEQUENTIAL` access pattern on our data as a whole. That way, each interface, effectively, is implemented as a FIFO.

Each of the streams internally are 9-bit value streams; that way, we could allow for delimiting 'bytes', which marked the ends of chunks or the end of the buffer overall (represented internally as `ENDOFCHUNK` and `ENDOFFILE` macros).

The one exception to this was the output from the `SHA` unit, as its output was always a set of 32 8-bit values matched to the number of chunks coming out of `LZW`.

3.5.2 Rabin Fingerprinting

The Rabin fingerprint reads data from a specialized function which creates a stream from the input buffer to the Hardware function.

This finds the chunk in this stream and upon finding the chunk, it delineates the implementation with a special `END-OFCHUNK` or `ENDOFFILE` stream byte. This allows the later applications to figure out where the chunk boundary is. The stream is at the same time sent to both the `SHA` component and the `LZW` component which work parallelly to process this stream as explained below.

The dataflow model allows us to focus at a stream level and use an additional stream byte to indicate important boundaries within the stream.

3.5.3 SHA-256

It takes streamed input from CDC, collects 64 bytes of data and performs its computation on that sub-chunk. Due to the sequential nature of SHA, it collects 64 bytes for each sub-chunk and computes them in a pipelined fashion. Once it is done with one chunk, it streams out the 256 bit (32 byte) SHA value for that chunk. Due to the dataflow pragma, SHA doesn't wait for the entire chunk to be read in before starting its computation. It starts operating as soon as there is some data available in the input stream.

3.5.4 LZW

Given that the algorithm for compression itself involves, conceptually, one memory lookup, one memory write, and one read from the input stream for each incoming byte, a streaming and pipelined approach seemed most appropriate.

As it stands, we implement this operation using a small hash table and simplistic hash function which uses as its 'key' the conceptual address of the 'large-memory' LZW-table implementation (which would require 8k rows of 256 columns each, each of which would contain a 13-bit value).

For this, we elected to hash this key down to 10 bits, using an arbitrary (but relatively effective) hash function to cut down on computational requirements. This meant we had 1024 hash rows, and we elected to allow for 24 entries into each; a factor of 8 was used to allow for the potential 8k table rows, and an extra factor of 3 was to allow for hash collisions and similar overflows.

A more careful implementation would allow for a small associative memory to handle such overflow; however, given the time constraints, we did not do this.

The one last consideration we had was the burden of 'resetting' the table for each chunk. Writing all of the values to our chosen 'Not Found' value (`0x1FFF`) would be a huge memory strain.

As such, we elected to have another table of validity bits, to mark which parts of the LZW hash table were valid for that iteration. This would require only 1024 24-bit values, which were easier to reset for each chunk than the entire hash table.

3.5.5 Deduplication

This function includes the calls to the `chunkDict` location of shared memory, containing the records of each `SHA` digest and its matching index.

Given the potentially large number of output chunks, we could not keep these records in the PL memory. As such, we implemented a rough dictionary in shared memory, hashing the digest down to 16 bits with a simple `XOR` strategy, and using that to index into our shared memory area.

We decided to allow for some degree of hash collisions here; with a potential worst-case of roughly 2^{17} chunks to store digests for, we allowed a factor-of-2 safety on our hash function, leading to 4 records per hash bucket. There's a reasonable

chance that this decision is leading to our woes of things not working; a larger stretch of time would have allowed for better analysis on this point.

The rest of the function was pedestrian by comparison; it reads the entire LZW stream (so no excess data is left) into a buffer, as well as the entire incoming SHA digest that corresponds to it.

Depending on whether or not there was a match in our dictionary, the relevant header is made for the output chunk, and either just that header, or the entire LZW-compressed section, is sent to the output stream.

3.6 Mapping to Zynq

Utilization Estimates					
Summary					
Name	BRAM_18K	DSP48E	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	506	-
FIFO	-	-	-	-	-
Instance	-	-	0	33	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	812	-
Register	0	-	988	32	-
Total	0	0	988	1383	0
Available	432	360	141120	70560	0
Utilization (%)	0	0	~0	1	0

Figure 2: CDC Mapping.

Utilization Estimates					
Summary					
Name	BRAM_18K	DSP48E	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	49	-
FIFO	-	-	-	-	-
Instance	3	1	2643	3242	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	66	-
Register	-	-	62	-	-
Total	3	1	2705	3357	0
Available	432	360	141120	70560	0
Utilization (%)	~0	~0	1	4	0

Figure 3: SHA Mapping.

Utilization Estimates					
Summary					
Name	BRAM_18K	DSP48E	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	88	-
FIFO	-	-	-	-	-
Instance	51	-	1687	2815	-
Memory	-	-	-	-	-
Multiplexer	-	-	-	96	-
Register	-	-	72	-	-
Total	51	0	1759	2999	0
Available	432	360	141120	70560	0
Utilization (%)	11	0	1	4	0

Figure 4: LZW Mapping.

Utilization Estimates

Summary

Name	BRAM_18K	DSP48E	FF	LUT	URAM
DSP	-	-	-	-	-
Expression	-	-	0	491	-
FIFO	-	-	-	-	-
Instance	-	-	8610	14929	-
Memory	5	-	16	4	-
Multiplexer	-	-	-	452	-
Register	-	-	494	-	-
Total	5	0	9120	15876	0
Available	432	360	141120	70560	0
Utilization (%)	1	0	6	22	0

Detail

Figure 5: Deduplicate Mapping.

3.7 Memory and computation model

Memory Requirements

(i) Content-Defined Chunking:

We'll need a rolling hash window's worth of working memory, spanning 16ish bytes. Structure to store rabin fingerprint, chunk size that has been slid over. 2 2KB tables that is used by the Rabin algorithm.

(ii) SHA-256:

Eight 32-bit span of memory to hold hash values, sixty four 32-bit span of memory for storing message schedule, sixty four 32-bit span of memory for storing round constants, ten 32-bit span of memory for working variables, and 64-byte SHA-block span of memory.

(iii) Chunk Matching:

We'll want a table to store hash values for index purposes, which would require at least 8 bytes times the maximum number of chunks to be processed.

(iv) LZW Encoding:

This is a somewhat tricky question given the associative memory involved, but it will be on the scale of roughly MAX_CHUNK.SIZE entries times 12 bits.

Computational Requirements

(i) Content-Defined Chunking:

Ignoring 1 time setup operations, then for each byte: rabinslide : 1 xor, 1 add, 1 mod rabinappend : 1 (OR + XOR + right shift + left shift) rabinnextchunk: 2 compares, 1 bitwise AND, 1AND, 1 OR

(ii) SHA-256:

Computation work per chunk: (Ignoring the index and loop iterator computations) Prepare the message schedule $m[i]$: $16 * (3 \text{ ORs} + 3 \text{ Shifts}) = 96$ operations $(64 - 16) * (3 \text{ ADDS} + 11 \text{ for SIG0} + 11 \text{ for SIG1}) = 1200$ operations Update the working variables: $64 * (7 \text{ adds} + 14 \text{ for EP0} + 14 \text{ for EP1} + 4 \text{ for CH} + 5 \text{ for MAJ}) = 2496$ operations Update hash values: 8 adds Total no of operations = $96 + 1200 + 2496 + 8 = 3800$ computations

(iii) Chunk Matching:

The only real operation here is a dictionary lookup, so computation can sensibly be considered negligible.

(iv) LZW Encoding:

With sensible dictionary lookup, there should be on the scale of one comparison operation per incoming byte of input, plus possibly a couple of additions as we loop through the incoming data.

Memory Access Requirements

(i) Content-Defined Chunking:

1 read for byte from buffer 1 read for the rabin window(for rolling hash) 1 write to the rabin window(for rolling hash) 1 read for mod table(rabin logic)

(ii) SHA-256:

Memory operations per chunk(Ignoring local - BRAM reads and writes): Prepare the message schedule $m[i]$: $16 * (4 \text{ reads}) = 64$ Initiate the working variables: 8 reads Update hash values : 8 writes Total operations: $64 + 8 + 8 = 80$

(iii) **Chunk Matching:**

One dictionary lookup should be required for each incoming hashed value; as such, we're looking at roughly 32 bytes of memory read (for reading the hash), and whatever memory costs are required after that for a dictionary lookup on that value.

(iv) **LZW Encoding:**

With efficient encoding, there should be roughly one memory read and one memory write involved in devising the code for each incoming byte as part of LZW. This will, of course, be very dependent on the specific dictionary implementation.

3.8 Area estimations

Unit	Unit cost	Area
Fixed Area	1	10
ARM Cortex-A53	1	2
ARM L2 Cache	1	$2.18(\text{data}) + 0.15(\text{tag}) = 2.33$
LUTs	$56272(\text{SHA}) + 3254(\text{LZW}) + 8285(\text{CDC})$	2.03
DSPs	$1(\text{SHA}) + 0(\text{LZW}) + 0(\text{CDC})$	0.01
Multipliers	$0(\text{SHA}) + 0(\text{CDC}) + 0(\text{LZW})$	0
DMA	$1(\text{SHA}) + 1(\text{LZW}) + 1(\text{SHA})$	0.3
BRAMs		$0.0092(\text{SHA}) + 1.37(\text{CDC}) + 10.1(\text{LZW}) = 11.47$

3.9 Bottleneck

Thanks to our decision to leverage the **DATAFLOW** pragma to handle our data movement in the PL, we have exactly zero idea of what the bottleneck in our design is.

I know, that's not a great answer. When we try to make traces of our program execution, no tools that xilinx provides us allow for any kind of per-element trace of execution, or logs of streams being full/empty, or any other kind of relevant information. The best we have to go on are estimates of latencies of our functions, but, given that some of them run across variable-length chunks a variable number of times within our hardware buffer, the ranges are huge, and not reflective of actual execution.

However, what we **can** do is make estimates of what the bottlenecks might be. For instance: with a hardware buffer a couple factors of scale past our largest chunk workload, we can estimate that the data movement to and from the FPGA is not the primary concern; that overhead should amortize out far past any dominating bottleneck term.

So where, in our core logic inside the PL, could the bottleneck be? It is possible that it is the compute within the **SHA** function; there are certainly a lot of loops to perform for each block of 512 bytes, and that necessarily takes time.

Alternatively, the **LZW** unit could be the bottleneck. Though the inner function, which operates on a per-chunk basis, has been pipelined down to an **II** of 6 (at 200MHz) or 9 (at 300MHz), the overhead of clearing the validity bits for its hash table could be styming its ability to efficiently handle data at a high throughput.

Our current bottleneck suspect is the access to shared memory for looking at the dictionary of **SHA** digests; this necessarily stalls the **deduplicate** unit until it can read a (current) total of 144 bytes from shared memory for each chunk, which it then compares against our incoming **SHA** value. Depending on how efficiently this operation can work, this may hang up the entire operation. It is also possible that having to wait for this to be done before beginning to output our constructed packet to our final output stream could be holding back the design.

The rest of our operations should be safe from bottlenecks; the reading of hardware buffer input into streams, the output from the final stream, and the process of Rabin-fingerprinting have all been pipelined effectively to initiation intervals of less than or equal to 2.

Obviously, there are a lot of potential culprits, but the current front-runners are likely the **SHA** unit and the lookup section of the **Deduplicate** unit.

4 10 Gbps design

As has been noted, the fact that we are operating on a **DATAFLOW** model means that our ability to peek inside the specifics of our design is limited. Xilinx does not like putting that many (read: any) calipers on the internals of our hardware function.

As such, we must examine possibilities for improvements in the abstract. In this section, we will discuss strategies for improving performance which, were we to work as diligently across the next 5 weeks as we have for this previous 5, would certainly yield performance improvements in pursuit of a *10Gbps* goal.

4.1 With Hardware Bounds in Mind

Looking at our current design, we can approach designing a 10Gbps design with the idea that the core functionality would still use logic almost entirely on the FPGA, but we could use more of our resources to solve the problem, and use as many tricks as possible to maximise the utility of the PL's time.

One of the first things we can see from our current resource usage is that we have a good number of BRAM's still available, as well as a large proportion of our FF/LUT resources. This opens us up to putting more SHA units (for the computational resources) and LZW units (for the memory resources) onto the PL, with the idea that one of the two (likely the former) is the bottleneck culprit.

With the idea that a singular `deduplicate` unit would still need to control output, we could link the `rabin` unit to the `deduplicate` unit by way of multiple parallel SHA units; for this example, we will imagine 2 such units. This would allow us to alternate which unit hashed which series of chunks, allowing for an idealized 2x speedup on that axis.

As for the LZW unit, we could easily fit another onto the PL in the same fashion.

If the resources are not the bottleneck, then we would need to look into the process by which we are transferring data. the use of `hls::stream` FIFO units has been handy for automatically handling a lot of timing requirements, but they carry the disadvantage of limiting the rate by which we are (conceptually) pipelining data through our application.

One way we could improve upon this would be to use multiple streams in parallel for our application; conceptually, this would be similar to, instead of sending water through a single garden hose, sending water through two parallel garden hoses.

If our bottleneck in the PL is access to the shared memory inside which our SHA dictionary resides, we could mitigate that problem by decreasing the amount of memory we need to read. We could do this by expanding our shared memory hash table region; if we're expecting something along the lines of 2^{16} unique chunks in a file, we could allocate 2^{19} , or more, total hash lines such that we would probabilistically only need to pull one hash line for any given SHA digest.

In this way, our total memory access would be in the range of 36 bytes per chunk, rather than our current 144. The disadvantage would be, we may need a particularly significant field of hash values to near-guarantee that we would not encounter any hash collisions that may break our design.

On our software side, we are currently using one thread/core purely to read in input, one to write to our output memory, and then one thread is copying memory to the FPGA working region, invoking the PL function, and then copying that memory out to a location from which the final thread writes.

A different implementation would more cleverly synchronize the operations that the middle thread performs, and make sure that the relevant memory copying did not need to happen (or at least, did not need to happen in a redundant fashion.)

4.2 With No Hardware Bounds in Mind

If we distance ourselves from the bounds of the Ultra96 system, however, we can consider all manner of other design options.

Let us approach this by keeping the core of our design (the use of streaming data, the `DATAFLOW` pragma, and doing LZW and SHA in parallel), and seeing what additional resources could add to our design.

One thing that could immediately help our design would be to have a fully-associative memory for the records of what the SHA digest was for each chunk index. Putting such a memory into small, local memories would allow for extremely fast lookups in service of deduplication.

By removing the need for shared memory access, that bottleneck could be completely eliminated. Further, it would only take a roughly 5MB memory to do so.

Considering our PL resources, if SHA is the bottleneck, the operation could be distributed across a huge number of computational units. In a similar fashion, given that each chunk is independent from the perspective of the LZW functionality, as many hardware units as possible could be thrown at that problem.

Going further, and expanding on the idea of parallel streams, any issues of transferring the data across FIFO units could be distributed to allow for, for instance, the `Rabin` computation or `Deduplication` units to pipeline data through at a quicker speed than would be otherwise possible, writing a larger number of bytes per FPGA "cycle" and further reduce the effective II of our `DATAFLOW`-pipelined design.

On the software side, we wouldn't anticipate needing sweeping structural changes. Instead, we would benefit from those cores being slightly faster. At 10Gbps, a new byte is coming in roughly every cycle for a 1.2GHz processor. With a 2.4GHz core, along with similarly-increased memory access capabilities, we could better assure that we did not drop packets somewhere in the process of moving all that data around.

As tempting as it would be to throw more threads at the task of reading input in, we fear that the overhead of synchronizing all such resources would outstrip the benefit of more processors doing the work.

That said, a dedicated data mover and hardware buffer to handle the incoming data could allow for the best and safest resource handling at 10Gbps.

5 Validation techniques

As of now the main validation technique is to run a script that can decode the compress.dat file generated from the application and then compares it with the input file that was sent over the network.

We also use prints on the main arm core, that prints the bytes read by the network thread, to provide real time testing validation that all bytes are being read and no packets are being dropped.

Since we are using separate cores for reading from network and writing to the file, it should not affect the timing of the core doing the processing. We are also using sync primitives to ensure correct order of operations.

6 Key lessons learned

- It is always good to define the interfaces first and make a working implementation first and then incrementally add functionalities so that we always have something to fall back to and start with a working code again.
- In our particular case of putting all pieces into hardware with streaming interfaces, this would have likely meant almost completely ignoring one of the milestones in an effort to put all of our separate pieces into hardware simultaneously. While this may not have made the grader happy, it would have pushed our hardware integration much earlier and allowed for less late-game hand-wringing.
- We learned how to make decisions when there is a memory-time trade off as there are limited memory and processing resources on the FPGA.
- We learned to use many more optimization techniques by using various pragmas to decrease the time it takes to run and also the memory it takes.
- While the DATAFLOW pragma allowed for a lot of data movement simplifications, Xilinx's stubborn determination to hide any reasonable or analysis information about its implementation ended up shooting us in the foot for both debugging and this report. In the future, we would pursue options that have mature debugging and analysis outputs.
- Careful debug output on our end through the process might have prevented some of the frustrations with Xilinx's systematic failure to elaborate on build errors.

7 Design space exploration

Describe design space explored and show graphs and models to support design selection For our earlier implementation in which we were running only LZW in hardware, the data was being sent to hardware for LZW and then the output back to the processor region to do the next steps in the computation, which led to very low speed of the overall implementation because of this data transfer overhead and most of the part still running in software.

But that incremental way of putting things in hardware helped us ultimately moving the entire design to the hardware. So, now all the components CDC, SHA, LZW and deduplication all run in hardware. This removed the delay of transferring intermediate results between the FPGA and the processor region and also increased the overall speed because of hardware acceleration techniques like using pipelining, unrolling, partitioning arrays etc.

To make advantage of running in hardware. we used dataflow pragma and serialized everything so that as soon as the data is available, the hardware component can start processing on it.

CDC provides chunk data to SHA and LZW simultaneously, which operate in parallel. The flow of our design is described in Figure 1.

To save space on FPGA memory, we use shared memory to store our hash table for the deduplication step. Another optimization technique was to use separate arm cores to read the input data from ethernet, to process on the data and to write the data to the output file.

That saves time in the sense that it doesn't have to wait for the data transmission to complete before it can start processing




Figure 6: Data flow.

on the data already read. So, one core reads the data from the ethernet and writes it to the shared memory, another core reads it and transfers it to the FPGA as soon as 2MB data has been written to it, and the third core writes the output data as soon as the 2MB data has been processed in the FPGA. This makes the overall implementation faster.

The changes in what speed we get while moving from all software to all hardware is shown in the following table:

Implementation	All Software	LZW in hardware	All hardware	All hardware with dataflow
Speed	7 Mbps	15 Mbps	10 Mbps	193 Mbps

Putting LZW in hardware increased the speed to 15 Mbps from 7 Mbps when everything was in software. After putting everything in hardware but still moving the data from software to hardware and reading it back from hardware to software for each component, it actually decreased the speed because of data transfer overhead. Also, deduplication was still running in software, so data flow from hardware to software for it added to the delays.

The new approach using dataflow and every component in hardware boosted the speed upto 193 Mbps for some large files for the reasons described above.

8 Individual contribution

Taylor implemented the LZW, chunkdict associative map and the basic framework for the Dataflow model. Nishanth worked on the Rabin CDC functionality and also implemented the network and output threads on each core.

Ritika worked on SHA software and hardware implementations.

All of us were involved in various capacities in the debug and testing phases of this project.

9 Academic code of integrity

We, Ritika, Taylor and Nishanth, certify that we have complied with the University of Pennsylvania's Code of Academic Integrity in completing this final exercise.